

Data

The data is **not stored in this repository**. Run the notebook to scrape it; files are written to `swiggy_project_outputs/`.

Source

Item	Value
Platform	Google Play Store (India)
App	Swiggy: Food Instamart Dineout (<code>in.swiggy.android</code>)
Language filter	English (<code>lang="en"</code> , <code>country="in"</code>)
Sort order	Newest first
Reviews scraped	3,000
Period in the reported run	23 Sep 2026 - 30 Sep 2026
Reviews after cleaning	1,145
Tool	<code>google-play-scraper</code>

Why the data is not committed

- Reviews are user-generated content owned by their authors and subject to Google Play's terms.
- Each run fetches the newest reviews, so a committed snapshot would go stale immediately.

Data dictionary: raw file (`01_swiggy_reviews_raw.csv`)

Column	Type	Description
<code>reviewId</code>	string	Unique review ID from Google Play
<code>content</code>	string	Review text
<code>score</code>	int (1-5)	Star rating given by the reviewer
<code>thumbsUpCount</code>	int	"Helpful" votes from other users
<code>reviewCreatedVersion</code>	string	App version when the review was first written
<code>at</code>	datetime	Review date
<code>replyContent</code>	string	Swiggy's public reply, if any
<code>repliedAt</code>	datetime	Reply date
<code>appVersion</code>	string	App version at the time of the review

`userName` and `userImage` are **dropped immediately after scraping**; they are personal data the analysis does not need.

Main derived columns (`08_swiggy_reviews_scored.csv`)

Column	Description
<code>review</code> , <code>rating</code>	Renamed <code>content</code> , <code>score</code>
<code>sentiment_text</code>	Lightly cleaned text for sentiment models (keeps case, punctuation, emojis)
<code>clean_text</code>	Fully cleaned, lemmatised text for TF-IDF and n-grams
<code>word_count</code> , <code>english_ratio</code>	Used by the length and language filters
<code>star_sentiment</code>	Label from stars: 1-2 Negative, 3 Neutral, 4-5 Positive
<code>vader_compound</code> , <code>vader_sentiment</code>	VADER score and label
<code>textblob_polarity</code> , <code>textblob_subjectivity</code> , <code>textblob_sentiment</code>	TextBlob scores and label
<code>roberta_neg</code> , <code>roberta_neu</code> , <code>roberta_pos</code> , <code>roberta_score</code> , <code>roberta_sentiment</code>	RoBERTa class probabilities, score and label
<code>literal_score</code> , <code>literal_sentiment</code>	Output of the best model before sarcasm correction
<code>complaint_context</code>	Review mentions a real problem (used by signal S1)
<code>sarcasm_cue</code>	Sarcasm cue phrase found (signal S2)
<code>irony_prob</code>	Irony-model probability (signal S3)
<code>is_sarcastic</code>	Any of the three sarcasm signals fired
<code>possible_rating_error</code>	Fully positive text with 1-2★ and no complaint: likely a mis-tap, not sarcasm

Column	Description
has_reply	Swiggy replied publicly
final_sentiment	Sarcasm-adjusted sentiment used in all later analysis
asp_*	One flag per business aspect mentioned in the review